

# Maximum Likelihood Estimation - Linear Case Exercise

HY-673, 2023-2024, Spring

## Exercise 3: MLE - Linear case (2022-2023 Assignment #1)

- (a) Let  $x_1, \dots, x_n \in \mathbb{R}$  be a given dataset. Assume the parametric model that generates the samples to be  $p_\theta(x) \equiv \mathcal{N}(0, \theta)$ , where  $\theta$  corresponds to the variance  $\sigma^2$ . Compute analytically the MLE solution  $\hat{\theta}_{\text{MLE}}$ .

### Solution:

Letting  $\mathbf{X} = x_1, x_2, \dots, x_n$ ,  $x_i \in \mathbb{R}$ ,  $i = 1, 2, \dots, n$ , we need to find by definition:

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p_{\theta}(\mathbf{X}|\theta), \quad (1)$$

and since we are given that  $p_{\theta}(x)$  is a univariate Gaussian, the likelihood function is:

$$\mathfrak{L} = p_{\theta}(\mathbf{X}|\theta) = \mathcal{N}(\mathbf{X}|\theta) = p_{\theta}(\mathbf{X}|\mu, \sigma^2), \quad \theta = \mu, \sigma^2. \quad (2)$$

Due to the Gaussian having an exponential, it will be easier to work with the logarithmic likelihood function which will be:

$$\ln \mathfrak{L} = \ln(\mathcal{N}(\mathbf{X}|\mu, \sigma^2)) = \sum_{n=1}^N \ln(\mathcal{N}(x_n|\mu, \sigma^2)) = \sum_{n=1}^N \ln\left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{(x_n-\mu)^2}{\sigma^2}\right)}\right) \quad (3)$$

$$= \sum_{n=1}^N -\ln\left(\sqrt{2\pi\sigma^2}\right) - \frac{(x_n - \mu)^2}{2\sigma^2} = -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2. \quad (4)$$

Since  $\ln(x) : \mathbb{R}^+ \rightarrow \mathbb{R}$  is a monotonically increasing function, we can equivalently find:

$$\hat{\mu}_{\text{MLE}} = \arg \max_{\mu} \ln \mathfrak{L}, \quad \text{and} \quad (5)$$

$$\hat{\sigma}_{\text{MLE}}^2 = \arg \max_{\sigma^2} \ln \mathfrak{L}. \quad (6)$$

It is not required for us to find  $\hat{\mu}_{\text{MLE}}$ , but, let's do it anyway. For that, we need the critical points of the logarithmic likelihood function with respect to  $\mu$ :

$$\frac{\partial}{\partial \mu} \ln \mathfrak{L} = 0 \implies \frac{\partial}{\partial \mu} \left( -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 \right) = 0 \implies \quad (7)$$

$$\frac{1}{\sigma^2} \sum_{n=1}^N (x_n - \mu) = 0 \implies \sum_{n=1}^N x_n - N\mu = 0 \implies \mu = \frac{1}{N} \sum_{n=1}^N x_n, \quad (8)$$

therefore, since it is given that  $p_\theta(X) \equiv \mathcal{N}(\mu = 0, \sigma^2 = \theta)$ :

$$\hat{\mu}_{\text{MLE}} = \frac{1}{N} \sum_{n=1}^N x_n = 0. \quad (9)$$

Likewise, to find  $\hat{\sigma}_{\text{MLE}}^2$ , we want:

$$\frac{\partial}{\partial \sigma^2} \ln \mathfrak{L} = 0 \implies \frac{\partial}{\partial \sigma^2} \left( -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 \right) = 0 \implies \quad (10)$$

$$-\frac{N}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{n=1}^N (x_n - \mu)^2 = 0 \implies -N\sigma^2 + \sum_{n=1}^N (x_n - \mu)^2 = 0 \implies \sigma^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu)^2, \quad (11)$$

hence, since  $p_\theta(X) \equiv \mathcal{N}(\mu = 0, \sigma^2 = \theta)$ , the final answer is:

$$\boxed{\hat{\theta}_{\text{MLE}} = \hat{\sigma}_{\text{MLE}}^2 = \frac{1}{N} \sum_{n=1}^N x_n^2.} \quad (12)$$

- (b) Generate samples  $x_i \sim \mathcal{N}(0, \theta^*)$  with  $\theta^* = 2$  and create one realization of the dataset. For  $n = 10^1, 10^2, 10^3$ , plot the log-likelihood as a function of  $\theta$ . What do you observe?

### Solution:

The requested code for the plot is in `03e_MLE_linear_case.ipynb`:




Figure 1: We observe that as  $n$  increases, it becomes easier for the log-likelihood functions reach the ground truth value  $\theta^* = 2$ .

- (c) Repeat (b) averaging over 100 realizations with new data at each realization and plot the average log-likelihood as a function of  $\theta$ . For each realization, compute also  $\hat{\theta}_{\text{MLE}}(n)$  and plot the respective histogram. Comment on the shape of the average log-likelihood function and the shape of the histogram of  $\hat{\theta}_{\text{MLE}}(n)$ .

## Solution:

The requested plots are below, and the code in `03e_MLE_linear_case.ipynb`.




Figure 2: The shape of the average log-likelihood functions do not differ much from the shape of the log-likelihood functions that we observed in the previous question. Taking the average log-likelihood function over 100 realizations of a Gaussian dataset helps significantly approximate  $\theta^* = 2$ , even with a small dataset comprised of  $n = 10$  samples. The histogram plot shows that by drawing more samples, our estimators become more “confident”, i.e., their variance drops. The dispersion of the predictions is lower for a bigger value of  $n$ , which follows the intuition that more data points result in greater estimation confidence.

(d) Compute analytically the Fisher information:

$$\mathcal{I}(\theta) := \mathbb{E}_{p_\theta} \left[ \left( \frac{d}{d\theta} \log p_\theta(x) \right)^2 \right]. \quad (13)$$

## Solution:

Setting  $p_\theta(X) \equiv \mathcal{N}(\mu = 0, \sigma^2 = \theta)$ , then from eq. [3,4,10,11](#), the Fisher information is:

$$\mathcal{I}(\theta) = -\mathbb{E}_{p_\theta} \left[ \frac{d^2}{d\theta^2} \ln p_\theta(X) \right] = -\mathbb{E}_{p_\theta} \left[ \frac{d}{d\theta} \left( -\frac{1}{2\theta} + \frac{1}{2\theta^2} (X - \mu)^2 \right) \right] \quad (14)$$

$$= -\mathbb{E}_{p_\theta} \left[ \frac{1}{2\theta^2} - \frac{1}{\theta^3} (X - \mu)^2 \right] = \frac{1}{2\theta^3} \mathbb{E}_{p_\theta} [(X - \mu)^2] = \frac{1}{2\theta^2}. \quad (15)$$